

SentinelX

3. Human-in-the-Loop Decision

One of the most important responsible AI decisions within SentinelX is the decision the system intentionally does NOT make — and the fact that the “human gate” looks different depending on the flow.

SentinelX does not autonomously determine whether a case should be escalated to intervention organizations, NGOs, law enforcement agencies, or emergency responders, in any of its three flows:

- For SOS, there is no AI decision to make — the alert is sent directly, and what happens next is entirely up to the receiving humans (emergency contacts and the NGO analyst).
- For Detect Threat, the human control point is the user's own consent: the on-device Bi-LSTM only ever runs while the user has deliberately toggled the feature on, and it never does anything beyond flagging a transcript snippet for an analyst to see.
- For Safe Chat + Upload Evidence, the human control point is the analyst's explicit [Analyze Report] click — Gemini never analyzes evidence until a human specifically asks it to.

This limitation was intentionally designed into the architecture because exploitation, abuse, coercion, stalking, and contextual threat cases are highly sensitive human situations that often require emotional interpretation, organizational judgment, contextual understanding, and ethical consideration beyond algorithmic inference.

A false escalation could create panic, damage trust, misclassify vulnerable individuals, or trigger inappropriate organizational responses. At the same time, a missed escalation could leave vulnerable individuals unsupported during dangerous situations.

Because of these real-world consequences, human analysts remain fully responsible for:

- reviewing all SOS alerts, Detect Threat flags, and Gemini evidence analyses,
- validating AI-generated insights rather than treating them as conclusions,
- interpreting narrative and evidentiary nuance,
- assessing intervention necessity,

- documenting the rationale behind any decision made,
- and making all final decisions regarding escalation or response.

The AI system supports organizational reasoning but never replaces human accountability.